# Competitive Landscape

HyperClaw enters a rapidly growing AI orchestration market against open-source frameworks that were built for research, not enterprise. We are the first platform designed for production-grade multi-agent workflows from day one.

## FEATURE MATRIX

| CAPABILITY | HYPERCLAW | LANGCHAIN | AUTOGEN | CREWAI |
|---|---|---|---|---|
| Intelligent Agent Routing | UCB1 Bandit WIN | Manual chains | Round-robin | Role-based |
| Memory Architecture | HyperMemory Causal Graph WIN | Buffer / vector | Conversation buffer | ✗ Limited |
| Policy / Safety Engine | HyperShield WIN | ✗ | Basic guardrails | ✗ |
| Specialist Agents OOTB | 50+ agents (5 domains) WIN | Templates only | ~6 examples | Role templates |
| State Management | HyperState (Pydantic v2) | Partial | ✓ Basic | Partial |
| Enterprise Auth / SSO | ✓ Roadmap Q2 2026 | ✗ | ✗ | ✗ |
| Multi-Model Routing | ✓ Native | ✓ Via LLMs | ✓ Yes | Limited |
| Recursive / Meta-Agents | ✓ PROMETHEUS + SOLOMON WIN | ✗ | Nested agents | ✗ |
| Self-Expanding Agents | ✓ GENESIS Protocol WIN | ✗ | ✗ | ✗ |
| Civilization Knowledge Base | ✓ SOPs, Org Charts, JDs WIN | ✗ | ✗ | ✗ |

| CAPABILITY | HYPERCLAW | LANGCHAIN | AUTOGEN | CREWAI |
|---|---|---|---|---|
| Three Human Gates | ✓ Ethical, Auth, Uncertainty WIN | ✗ | Partial | ✗ |
| 30+ Platform Connectors | ✓ HyperClaw Gateway WIN | Via LangSmith | Limited | Limited |
| Managed Cloud Offering | ✓ Supabase-backed | LangSmith (paid) | ✗ | ✗ |
| Test Coverage | 377 passing WIN | Varies | Varies | Partial |
| License | MIT (open core) | MIT | MIT/CC | MIT |
| Enterprise SLA | ✓ Available | ✗ | ✗ | ✗ |

## OUR SIX UNFAIR ADVANTAGES

ROUTING

### UCB1 HyperRouter — Self-Improving

HyperRouter uses Upper Confidence Bound (UCB1) bandit algorithms to learn optimal model-to-task assignment. Gets smarter with every interaction. Formula: ucb1_score = mean_reward + C × √(ln(total) / attempts). No competitor has this.

MEMORY

### Three-Tier Civilizational Memory

Session, Domain, and Civilizational Graph (pgvector) — permanent, growing memory. Causal relationships preserved across tiers. LangChain's buffer and CrewAI's lack of persistent memory are no match.

SAFETY

### HyperShield + Three Human Gates

Zero-trust, hot-reload YAML policies. Three Human Gates baked in: Ethical, Authorization, and Uncertainty (<60% confidence escalates to human review). No competitor ships a dedicated policy engine.

AGENTS

### 50+ Specialist Agents OOTB

Across 5 domains, HyperClaw ships production-ready specialist agents managed by SOLOMON overmind. Competitors ship templates and examples — customers must build everything themselves.

SELF-EXPANSION

### GENESIS Protocol — Unique in Market

SOLOMON automatically detects capability gaps and creates new specialist agents via the GENESIS Protocol. The agent network grows with the

KNOWLEDGE

### Civilization Knowledge Base

Ingests SOPs, org charts, job descriptions, client profiles, runbooks, and policies. Procedural chunker preserves step order. Interview agent

organization. No other platform has self-expanding agents — period.

extracts tacit knowledge. OrgGraph traversal provides role-aware context injection to every agent.

## PRICING COMPARISON

### HYPERCLAW
**$499**
/ mo — Starter

- 50+ agents included
- HyperShield policy
- UCB1 routing
- Supabase cloud
- Enterprise SLA

### LANGCHAIN
**$0**
OSS + LangSmith $39+/mo

- Framework only
- No agents OOTB
- No policy engine
- LangSmith for obs.
- No enterprise SLA

### AUTOGEN
**$0**
OSS (Microsoft)

- Research focus
- ~6 agent examples
- Basic guardrails
- No cloud offering
- No SLA

### CREWAI
**$0**
OSS + paid tiers

- Role templates
- No memory graph
- No policy engine
- Enterprise: custom
- Growing ecosystem

## GITHUB STARS TRAJECTORY

### GitHub Stars (Approximate — as of Q1 2026)

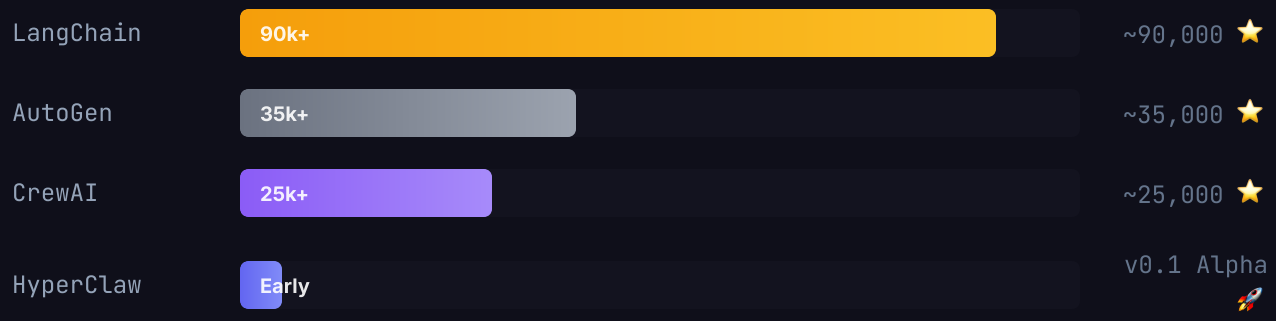


*HyperClaw launched in alpha (Q1 2026). LangChain, AutoGen, and CrewAI had multi-year head starts. HyperClaw's differentiation is enterprise architecture, not star count — the same position Snowflake held vs. Hadoop in 2012.*

## ENTERPRISE READINESS

### HyperClaw

| | |
|---|---|
| Policy Engine | ✓ |
| Audit Logging | ✓ |
| Enterprise SLA | ✓ |
| SSO (Q2 2026) | Roadmap |

### LangChain

| | |
|---|---|
| Policy Engine | ✗ |
| Audit Logging | LangSmith only |
| Enterprise SLA | ✗ |
| SSO | ✗ |

### AutoGen / CrewAI

| | |
|---|---|
| Policy Engine | ✗ |
| Audit Logging | ✗ |
| Enterprise SLA | ✗ |
| SSO | ✗ |

| | |
|---|---|
| Dedicated Support | ✓ |
| Custom Deployment | ✓ |

| | |
|---|---|
| Dedicated Support | Paid plans |
| Custom Deployment | DIY |

| | |
|---|---|
| Dedicated Support | ✗ |
| Custom Deployment | DIY |

The competitive landscape reveals a clear gap: every major framework was built for developers experimenting with agents, not enterprises deploying them at scale. HyperClaw fills this gap with a production-first architecture, enterprise SLAs, and a policy engine that compliance teams can actually trust.

---